

Partitioning & Hierarchical Clustering

Likhita Pula

6/4/2021

Importing data for analysis For the first assignment on Partitioning and hierarchical clustering I have chosen a dataset that describes various attributes of wine. This dataset is hosted on www.kaggle.com.

```
wine_data <- read.delim("C:\\Users\\likhi\\Desktop\\MSDA\\Data Mining 2\\Dataset\\wine-clustering.csv",  
                        header = TRUE,  
                        sep=",")
```

Exploratory data analysis As a next step, we can look at the overview and statistical details of the dataset used.

Below we can see the top 6 rows of the data set

```
head(wine_data)
```

```
##   Alcohol Malic_Acid  Ash Ash_Alcanity Magnesium Total_Phenols Flavanoids  
## 1   14.23      1.71 2.43      15.6      127      2.80      3.06  
## 2   13.20      1.78 2.14      11.2      100      2.65      2.76  
## 3   13.16      2.36 2.67      18.6      101      2.80      3.24  
## 4   14.37      1.95 2.50      16.8      113      3.85      3.49  
## 5   13.24      2.59 2.87      21.0      118      2.80      2.69  
## 6   14.20      1.76 2.45      15.2      112      3.27      3.39  
##   Nonflavanoid_Phenols Proanthocyanins Color_Intensity Hue OD280 Proline  
## 1           0.28           2.29           5.64 1.04 3.92 1065  
## 2           0.26           1.28           4.38 1.05 3.40 1050  
## 3           0.30           2.81           5.68 1.03 3.17 1185  
## 4           0.24           2.18           7.80 0.86 3.45 1480  
## 5           0.39           1.82           4.32 1.04 2.93 735  
## 6           0.34           1.97           6.75 1.05 2.85 1450
```

Checking if the data set has any missing values:

```
sum(is.na(wine_data))
```

```
## [1] 0
```

We can see that there are no null values in the data.

The wine dataset had below features:

- Alcohol

- Malic acid
- Ash
- Alkalinity of ash
- Magnesium
- Total phenols
- Flavanoids
- Nonflavanoid phenols
- Proanthocyanins
- Color intensity
- Hue
- OD280 of diluted wines
- Proline

Checking the dimension of the dataset:

```
dim(wine_data)
```

```
## [1] 178 13
```

we can see that there are a total of 178 rows and 13 columns (features) in the wine dataset.

Checking the statistical attributes like maximum, minimum, mean etc. of the dataset features.

```
summary(wine_data)
```

```
##      Alcohol      Malic_Acid      Ash      Ash_Alcanity
## Min.   :11.03  Min.   :0.740  Min.   :1.360  Min.   :10.60
## 1st Qu.:12.36  1st Qu.:1.603  1st Qu.:2.210  1st Qu.:17.20
## Median :13.05  Median :1.865  Median :2.360  Median :19.50
## Mean   :13.00  Mean   :2.336  Mean   :2.367  Mean   :19.49
## 3rd Qu.:13.68  3rd Qu.:3.083  3rd Qu.:2.558  3rd Qu.:21.50
## Max.   :14.83  Max.   :5.800  Max.   :3.230  Max.   :30.00
##      Magnesium      Total_Phenols      Flavanoids      Nonflavanoid_Phenols
## Min.   : 70.00  Min.   :0.980  Min.   :0.340  Min.   :0.1300
## 1st Qu.: 88.00  1st Qu.:1.742  1st Qu.:1.205  1st Qu.:0.2700
## Median : 98.00  Median :2.355  Median :2.135  Median :0.3400
## Mean   : 99.74  Mean   :2.295  Mean   :2.029  Mean   :0.3619
## 3rd Qu.:107.00  3rd Qu.:2.800  3rd Qu.:2.875  3rd Qu.:0.4375
## Max.   :162.00  Max.   :3.880  Max.   :5.080  Max.   :0.6600
##      Proanthocyanins      Color_Intensity      Hue      OD280
## Min.   :0.410  Min.   : 1.280  Min.   :0.4800  Min.   :1.270
## 1st Qu.:1.250  1st Qu.: 3.220  1st Qu.:0.7825  1st Qu.:1.938
## Median :1.555  Median : 4.690  Median :0.9650  Median :2.780
## Mean   :1.591  Mean   : 5.058  Mean   :0.9574  Mean   :2.612
## 3rd Qu.:1.950  3rd Qu.: 6.200  3rd Qu.:1.1200  3rd Qu.:3.170
## Max.   :3.580  Max.   :13.000  Max.   :1.7100  Max.   :4.000
##      Proline
## Min.   : 278.0
## 1st Qu.: 500.5
## Median : 673.5
## Mean   : 746.9
## 3rd Qu.: 985.0
## Max.   :1680.0
```

We can see that few attributes like “proline” or “Magnesium” have larger values compared to other attributes. These attributes / features with larger variance can substantially influence output clusters, thus it would be good to scale our dataset.


```
scale_M <- scale(  
  x = wine_data  
)
```

scaling the data for analysis

Partitioning Clustering

Finding the optimal number of clusters We can use the Elbow method (sum of squares) method to find the optimal number of clusters:


```
set.seed(813)  
factoextra::fviz_nbclust(  
  x = scale_M,  
  FUNcluster = kmeans,  
  method = "wss"  
)
```



We can see the elbow points at 2 & 3. So, probably $K = 3$ or 2 would be giving us ideal clusters. Further we can observe the silhouette plot to find the best k value.

To find the optimal number of clusters we can examine the silhouette plot below:

```
set.seed(813)
factoextra::fviz_nbclust(
  x = scale_M,
  FUNcluster = kmeans,
  method = "silhouette"
)
```




From the silhouette plot we can say that we can get ideal number of clusters when $K=3$.

Finding the optimal K (number of clusters) value using gap statistics:

```
set.seed(813)
clusGap_kmeans <- cluster::clusGap(
  x = scale_M,
  FUNcluster = kmeans,
  K.max = 10
)

factoextra::fviz_gap_stat(
  gap_stat = clusGap_kmeans,
)
```



From the gap statistics we can see that the optimal value of K is 3.

k-means clustering

k-means clustering aims to partition the points into k groups such that the sum of squares from points to the assigned cluster centers is minimized. It takes number of groups, or initial group centers, then labels observations into groups that are close to common group centers. Then recalculates the centers and repeats.

```
kmeans_M1 <- kmeans(
  x = scale_M,
  centers = 3
)
kmeans_M1
```

```
## K-means clustering with 3 clusters of sizes 62, 51, 65
```

```
##
```

```
## Cluster means:
```

	Alcohol	Malic_Acid	Ash	Ash_Alcanity	Magnesium	Total_Phenols
## 1	0.8328826	-0.3029551	0.3636801	-0.6084749	0.57596208	0.88274724
## 2	0.1644436	0.8690954	0.1863726	0.5228924	-0.07526047	-0.97657548
## 3	-0.9234669	-0.3929331	-0.4931257	0.1701220	-0.49032869	-0.07576891


	Flavanoids	Nonflavanoid_Phenols	Proanthocyanins	Color_Intensity	Hue
## 1	0.97506900	-0.56050853	0.57865427	0.1705823	0.4726504
## 2	-1.21182921	0.72402116	-0.77751312	0.9388902	-1.1615122
## 3	0.02075402	-0.03343924	0.05810161	-0.8993770	0.4605046

	OD280	Proline
## 1	0.7770551	1.1220202

```
## 2 -1.2887761 -0.4059428
## 3  0.2700025 -0.7517257
##
## Clustering vector:
##  [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## [38] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 3 3 2 3 3 3 3 3 3 3 3 3 1
## [75] 3 3 3 3 3 3 3 3 3 2 3 3 3 3 3 3 3 3 3 1 3 3 3 3 3 3 3 3 3 3 3 3
## [112] 3 3 3 3 3 3 3 2 3 3 1 3 3 3 3 3 3 3 3 2 2 2 2 2 2 2 2 2 2 2 2 2
## [149] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
##
## Within cluster sum of squares by cluster:
## [1] 385.6983 326.3537 558.6971
## (between_SS / total_SS =  44.8 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
## [6] "betweenss"    "size"         "iter"         "ifault"
```

When we observe the kmeans clustering vector, we can clearly see that the dataset has been well clustered into three groups and the cluster means of each feature tell that all features are substantially influencing the output clusters (due to scaling).

```
cluster::clusplot(
  scale_M,
  kmeans_M1$cluster,
  color=TRUE,
  shade=FALSE,
  lines=0
)
```



These two components explain 55.41 % of the point variability.


The clusters formed using `kmeans()` look almost perfect. The dataset is well divided into 3 clusters in the `custplot` of `kmeans`.

Partitioning using Clara method

The `cluster::clara()` function is a good partitioning method for large data when robustness is not needed.

```
clara_M <- cluster::clara(  
  x = scale_M,  
  k = 3  
)  
plot(clara_M)
```

clusplot(cluster::clara(x = scale_M, k = 3))



These two components explain 55.41 % of the point variability.

Silhouette plot of cluster::clara(x = scale_M, k = 3)

n = 46


3 clusters C_j

$j: n_j \mid \text{ave}_{i \in C_j} s_i$

1 : 18 | 0.37

2 : 13 | 0.19

3 : 15 | 0.33



Average silhouette width : 0.31

When we look at the clustplot obtained from clara, we can see that there are 3 clusters created and the clusters created look good as the boundaries of these clusters are able to partition the clusters well.

But there is a slight overlap between two of the clusters which appears to be more than the overlap that is observed in clusters made from kmeans.

Results:

```
print(clara_M)
```

```
## Call:      cluster::clara(x = scale_M, k = 3)
## Medoids:
##           Alcohol Malic_Acid      Ash Ash_Alcanity  Magnesium Total_Phenols
## [1,]  1.3542080 -0.2831754  0.12204803  -0.2080942  0.2281415    0.7268305
## [2,] -0.9246039 -0.5427655 -0.89856839  -0.1482061 -1.3822227   -1.0307762
## [3,]  0.3934117  0.8088930  0.04914686   0.6003946 -0.5420327   -0.5833854
##           Flavanoids Nonflavanoid_Phenols Proanthocyanins Color_Intensity
## [1,]  0.8917481025          -0.33630220      1.37868246      0.4925666
## [2,]  0.0007311716           0.06545479      0.06831575     -0.7152224
## [3,] -1.2707199546           0.70826598     -0.59560339      1.4501706
##           Hue      OD280      Proline
## [1,]  0.4924084  0.1948119  0.9942817
## [2,]  0.1861586  0.7863692 -0.7522631
## [3,] -1.7825902 -1.3967588 -0.3076880
## Objective function:  2.821054
## Clustering vector:   int [1:178] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...
## Cluster sizes:      69 60 49
```

```
## Best sample:
## [1] 3 7 11 13 17 20 24 30 36 37 39 47 49 50 51 55 56 62 71
## [20] 77 85 95 102 107 108 115 116 120 122 124 129 133 136 137 142 147 149 152
## [39] 154 155 157 159 161 163 164 177
##
## Available components:
## [1] "sample" "medoids" "i.med" "clustering" "objective"
## [6] "clusinfo" "diss" "call" "silinfo" "data"
```

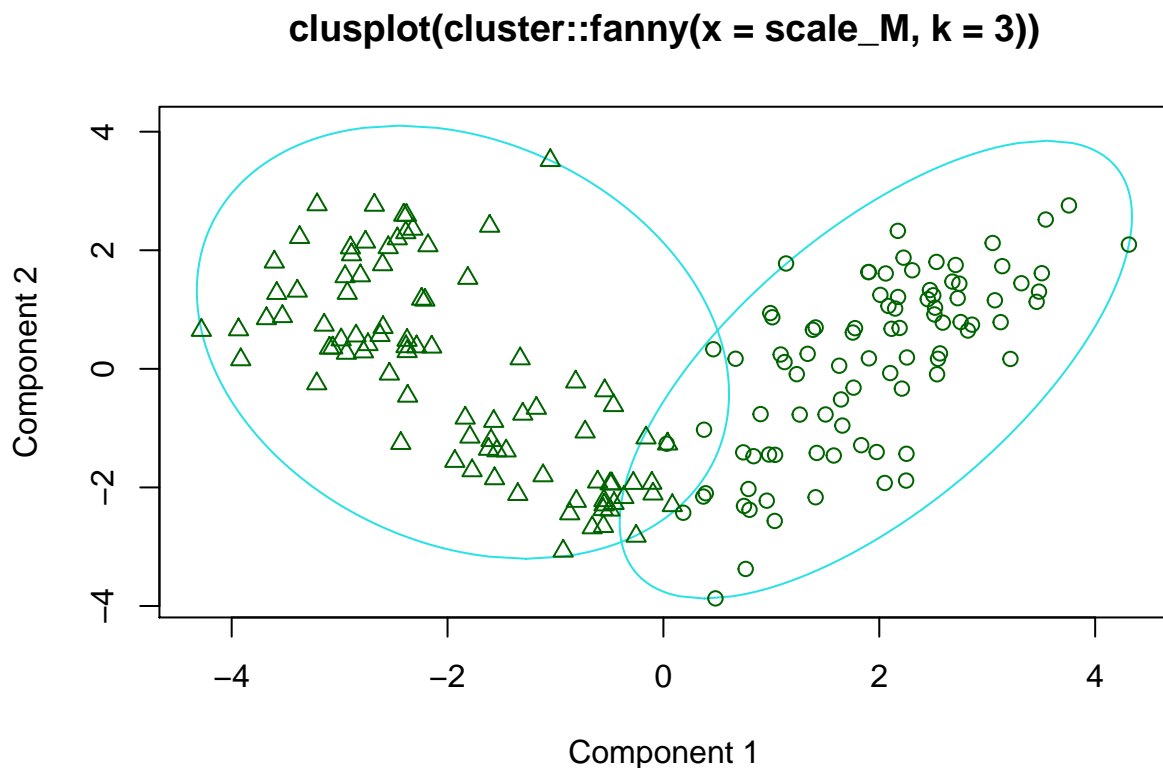
Partitioning using fanny method

The `cluster::fanny()` function gives a likelihood of a point belonging to a cluster.

```
fanny_M <- cluster::fanny(
  x = scale_M,
  k = 3
)
```

```
## Warning in cluster::fanny(x = scale_M, k = 3): the memberships are all very
## close to 1/k. Maybe decrease 'memb.exp' ?
```

```
plot(fanny_M)
```



These two components explain 55.41 % of the point variability.

Silhouette plot of cluster::fanny(x = scale_M, k = 3)


n = 178

2 clusters C_j

$j: n_j \mid \text{ave}_{i \in C_j} s_i$

1 : 89 | 0.28

2 : 89 | 0.24



Average silhouette width : 0.26

When we observe the clustplot created by fanny method, we can see that the two clusters formed are overlapping each other. However the silhouette plot looks good here suggesting that the two clusters are almost perfect. But looking at the overall picture we can say that clusters created by clara & kmeans look more promising than the clusters created using fanny.

Results:

```
print(fanny_M)
```

```
## Fuzzy Clustering object of class 'fanny' :
## m.ship.expon.      2
## objective      144.3284
## tolerance      1e-15
## iterations      37
## converged       1
## maxit          500
## n             178
## Membership coefficients (in %, rounded):
##      [,1] [,2] [,3]
## [1,]  33  33  33
## [2,]  33  33  33
## [3,]  33  33  33
## [4,]  33  33  33
## [5,]  33  33  33
## [6,]  33  33  33
## [7,]  33  33  33
```

##	[8,]	33	33	33
##	[9,]	33	33	33
##	[10,]	33	33	33
##	[11,]	33	33	33
##	[12,]	33	33	33
##	[13,]	33	33	33
##	[14,]	33	33	33
##	[15,]	33	33	33
##	[16,]	33	33	33
##	[17,]	33	33	33
##	[18,]	33	33	33
##	[19,]	33	33	33
##	[20,]	33	33	33
##	[21,]	33	33	33
##	[22,]	33	33	33
##	[23,]	33	33	33
##	[24,]	33	33	33
##	[25,]	33	33	33
##	[26,]	33	33	33
##	[27,]	33	33	33
##	[28,]	33	33	33
##	[29,]	33	33	33
##	[30,]	33	33	33
##	[31,]	33	33	33
##	[32,]	33	33	33
##	[33,]	33	33	33
##	[34,]	33	33	33
##	[35,]	33	33	33
##	[36,]	33	33	33
##	[37,]	33	33	33
##	[38,]	33	33	33
##	[39,]	33	33	33
##	[40,]	33	33	33
##	[41,]	33	33	33
##	[42,]	33	33	33
##	[43,]	33	33	33
##	[44,]	33	33	33
##	[45,]	33	33	33
##	[46,]	33	33	33
##	[47,]	33	33	33
##	[48,]	33	33	33
##	[49,]	33	33	33
##	[50,]	33	33	33
##	[51,]	33	33	33
##	[52,]	33	33	33
##	[53,]	33	33	33
##	[54,]	33	33	33
##	[55,]	33	33	33
##	[56,]	33	33	33
##	[57,]	33	33	33
##	[58,]	33	33	33
##	[59,]	33	33	33
##	[60,]	33	33	33
##	[61,]	33	33	33

##	[62,]	33	33	33
##	[63,]	33	33	33
##	[64,]	33	33	33
##	[65,]	33	33	33
##	[66,]	33	33	33
##	[67,]	33	33	33
##	[68,]	33	33	33
##	[69,]	33	33	33
##	[70,]	33	33	33
##	[71,]	33	33	33
##	[72,]	33	33	33
##	[73,]	33	33	33
##	[74,]	33	33	33
##	[75,]	33	33	33
##	[76,]	33	33	33
##	[77,]	33	33	33
##	[78,]	33	33	33
##	[79,]	33	33	33
##	[80,]	33	33	33
##	[81,]	33	33	33
##	[82,]	33	33	33
##	[83,]	33	33	33
##	[84,]	33	33	33
##	[85,]	33	33	33
##	[86,]	33	33	33
##	[87,]	33	33	33
##	[88,]	33	33	33
##	[89,]	33	33	33
##	[90,]	33	33	33
##	[91,]	33	33	33
##	[92,]	33	33	33
##	[93,]	33	33	33
##	[94,]	33	33	33
##	[95,]	33	33	33
##	[96,]	33	33	33
##	[97,]	33	33	33
##	[98,]	33	33	33
##	[99,]	33	33	33
##	[100,]	33	33	33
##	[101,]	33	33	33
##	[102,]	33	33	33
##	[103,]	33	33	33
##	[104,]	33	33	33
##	[105,]	33	33	33
##	[106,]	33	33	33
##	[107,]	33	33	33
##	[108,]	33	33	33
##	[109,]	33	33	33
##	[110,]	33	33	33
##	[111,]	33	33	33
##	[112,]	33	33	33
##	[113,]	33	33	33
##	[114,]	33	33	33
##	[115,]	33	33	33

##	[116,]	33	33	33
##	[117,]	33	33	33
##	[118,]	33	33	33
##	[119,]	33	33	33
##	[120,]	33	33	33
##	[121,]	33	33	33
##	[122,]	33	33	33
##	[123,]	33	33	33
##	[124,]	33	33	33
##	[125,]	33	33	33
##	[126,]	33	33	33
##	[127,]	33	33	33
##	[128,]	33	33	33
##	[129,]	33	33	33
##	[130,]	33	33	33
##	[131,]	33	33	33
##	[132,]	33	33	33
##	[133,]	33	33	33
##	[134,]	33	33	33
##	[135,]	33	33	33
##	[136,]	33	33	33
##	[137,]	33	33	33
##	[138,]	33	33	33
##	[139,]	33	33	33
##	[140,]	33	33	33
##	[141,]	33	33	33
##	[142,]	33	33	33
##	[143,]	33	33	33
##	[144,]	33	33	33
##	[145,]	33	33	33
##	[146,]	33	33	33
##	[147,]	33	33	33
##	[148,]	33	33	33
##	[149,]	33	33	33
##	[150,]	33	33	33
##	[151,]	33	33	33
##	[152,]	33	33	33
##	[153,]	33	33	33
##	[154,]	33	33	33
##	[155,]	33	33	33
##	[156,]	33	33	33
##	[157,]	33	33	33
##	[158,]	33	33	33
##	[159,]	33	33	33
##	[160,]	33	33	33
##	[161,]	33	33	33
##	[162,]	33	33	33
##	[163,]	33	33	33
##	[164,]	33	33	33
##	[165,]	33	33	33
##	[166,]	33	33	33
##	[167,]	33	33	33
##	[168,]	33	33	33
##	[169,]	33	33	33


[illegible]

Partitioning using pam method

The cluster::pam() algorithm is the robust version of k-means. It uses medoids, and centers are observations in the data set.

```
pam_M <- cluster::pam(
  x = scale_M,
  k = 3
)
plot(pam_M)
```

clusplot(cluster::pam(x = scale_M, k = 3))



These two components explain 55.41 % of the point variability.

Silhouette plot of cluster::pam(x = scale_M, k = 3)

n = 178


3 clusters C_j

$j: n_j \mid \text{ave}_{i \in C_j} s_i$

1 : 74 | 0.25

2 : 55 | 0.22

3 : 49 | 0.34



Average silhouette width : 0.27

We can see that there is significant overlap between the clusters formed by pam(). Kmeans, clara & fanny are giving better results compared to pam().

Results:

```
print(pam_M)
```

```
## Medoids:
##      ID      Alcohol Malic_Acid      Ash Ash_Alcanity  Magnesium
## [1,]  36  0.5904981 -0.4711544  0.15849862  0.3009543  0.01809398
## [2,] 107 -0.9246039 -0.5427655 -0.89856839 -0.1482061 -1.38222271
## [3,] 149  0.3934117  0.8088930  0.04914686  0.6003946 -0.54203270
##      Total_Phenols  Flavanoids Nonflavanoid_Phenols Proanthocyanins
## [1,]      0.6469393  0.9518166597      -0.81841060      0.47016154
## [2,]     -1.0307762  0.0007311716      0.06545479      0.06831575
## [3,]     -0.5833854 -1.2707199546      0.70826598     -0.59560339
##      Color_Intensity      Hue      OD280      Proline
## [1,]      0.01807806  0.3611585  1.2089101  0.5497067
## [2,]     -0.71522236  0.1861586  0.7863692 -0.7522631
## [3,]      1.45017064 -1.7825902 -1.3967588 -0.3076880
## Clustering vector:
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## [38] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 1 2 1 2 2 2 1 2 1 2 1
## [75] 1 2 2 2 2 1 2 2 2 3 2 2 2 2 2 2 2 2 2 2 1 1 2 1 2 2 2 2 2 2 2 2 1 1
## [112] 2 2 2 2 2 2 2 2 2 1 1 2 2 1 2 2 2 2 2 3 3 3 3 3 3 3 3 3 3 3 3 3 3
## [149] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
```

```
## Objective function:
##      build      swap
## 2.910808 2.806293
##
## Available components:
## [1] "medoids"      "id.med"      "clustering" "objective" "isolation"
## [6] "clusinfo"     "silinfo"     "diss"       "call"       "data"
```

Final Thoughts on Partitioning Algorithm below are the key observations from partitioning algorithm:

- The best value for number of clusters (K) is 3 which is seen from the plots created using elbow method, Silhouette calculations and gap statistics.
- The best model for partitioning algorithm is given by kmeans() as the clusters are well separated from each other.

Hierarchical Clustering

sample of distance matrix:

```
dist(
  x = scale_M[1:5,]
)
```


```
##           1           2           3           4
## 2 3.487697
## 3 3.018094 4.131258
## 4 2.834509 4.348349 3.237354
## 5 3.556821 4.614454 2.972721 4.483310
```

creating the distance matrix for hierarchical clustering:

```
dist_M <- dist(
  x = scale_M
)
```

Now that we have created our distance matrix, we can plot the qgraph that would show association between rows based on thickness of the lines.


```
library("qgraph")
qgraph::qgraph(
  input = 1/dist_M,
  layout="spring",
  minimum = 0.3
)
```



From the plot it looks like the dataset can be clustered into 3 groups. Further we can use `hclust()` function to execute agglomerative clustering which is joining the two most similar clusters and continuing until there is just a single cluster.

```
hclust_M <- hclust(
  d = dist_M,
  method = "complete"
)
plot(
  x = hclust_M
)
```

Cluster Dendrogram



The above dendrogram shows a complete linkage method where we can see that if we cut this dendrogram at a height of 8.8 (approx) we might get 3 clusters that are splitting the data into (approximately) equal sized clusters.

Lets cut the complete linkage tree so that we are able to get 3 clusters.

```
cutree_hclust_M <- cutree(  
  tree = hclust_M,  
  k = 3  
)
```

Plotting the clusters created by hierarchical clustering using complete linkage method:

```
require(ggplot2)
```


```
## Loading required package: ggplot2
```

```
prcomp_M <- data.frame(  
  prcomp(  
    x = scale_M,  
    center = FALSE,  
    scale. = FALSE  
  )$x[,1:2],  
  Cluster = as.character(cutree_hclust_M),  
  stringsAsFactors = FALSE  
)
```

```
ggplot(prcomp_M) +
  aes(x = PC1, y = PC2, color = Cluster, fill = Cluster, group = Cluster) +
  geom_point() +
  ggtitle("Complete Linkage Clustering", "Color corresponds to Hierarchical clusters") +
  theme_bw() +
  theme(legend.position = "none")
```

Complete Linkage Clustering

Color corresponds to Hierarchical clusters



Finding the coefficient of a hierarchical clustering using complete linkage method:

```
cluster::coef.hclust(hclust_M)
```

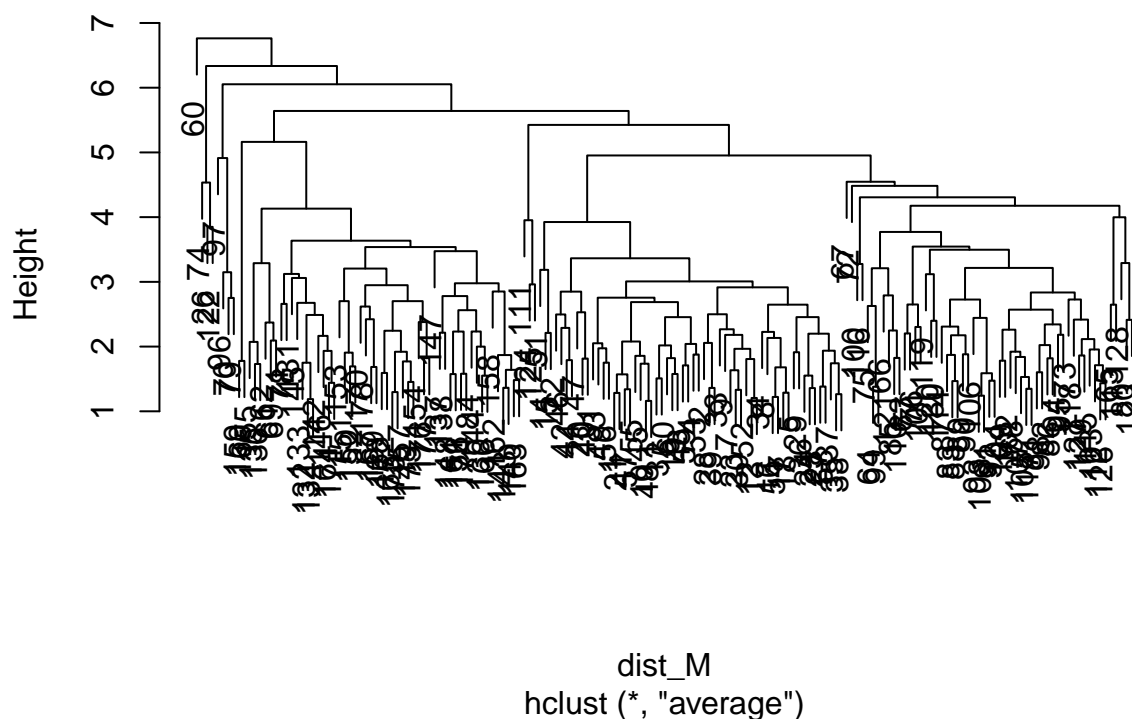
```
## [1] 0.815931
```

coefficient of 0.815 for hierarchical clustering using complete linkage method suggests that the hierarchical clustering is doing a better job of creating clusters with approximately equal number of observations. This is something we see from the cluster plot created earlier as well.

Further we can check the dendrogram for average linkage method:

```
hclust_average_M <- hclust(
  d = dist_M,
  method = "average"
)
plot(
  x = hclust_average_M
)
```

Cluster Dendrogram



From the average linkage dendrogram we can see that record 60 appears on a different branch of the tree. This could be an outlier in the wine data set.

Finding the coefficient of a hierarchical clustering using average linkage method:

```
cluster::coef.hclust(hclust_average_M)
```

```
## [1] 0.7006964
```

This coefficient helps us to detect outliers in our dataset but we can further check hierarchical clustering using single linkage.

Checking dendrogram for hierarchical clustering using single linkage method:

```
hclust_single_M <- hclust(
  d = dist_M,
  method = "single"
)
plot(
  x = hclust_single_M
)
```

Cluster Dendrogram



Finding the coefficient of a hierarchical clustering using single linkage method:

```
cluster::coef.hclust(hclust_single_M)
```

```
## [1] 0.5379128
```

We can see the coefficient for hierarchical clustering using single linkage is low (around 0.53) which indicates that this model is better for outlier detection.

Analysis of hierarchical clustering parameters To check the best hierarchical model for outlier detection and partitioning data into (approximately) equal sized groups I have created dendograms for different combinations of distance metric and methods of hierarchical clustering. (Which are key parameters of hierarchical clustering)

First, I am creating distance metric using “canberra”, “manhattan”, “euclidean”, “maximum” and “minkowski” method and similarly defining different hierarchical clustering methods like “ward.D”, “ward.D2”, “complete”, “mcquitty”, “average” and “single” linkage in `v_hclust`.

```
v_dist <- c(
  "canberra", "manhattan", "euclidean", "maximum", "minkowski"
)
list_dist <- lapply(
  X = v_dist,
  FUN = function(distance_method) dist(
    x = scale_M,
    method = distance_method
  )
)
```

```

)
)
names(list_dist) <- v_dist
v_hclust <- c(
  "ward.D", "ward.D2", "complete", "mcquitty", "average", "single"
)

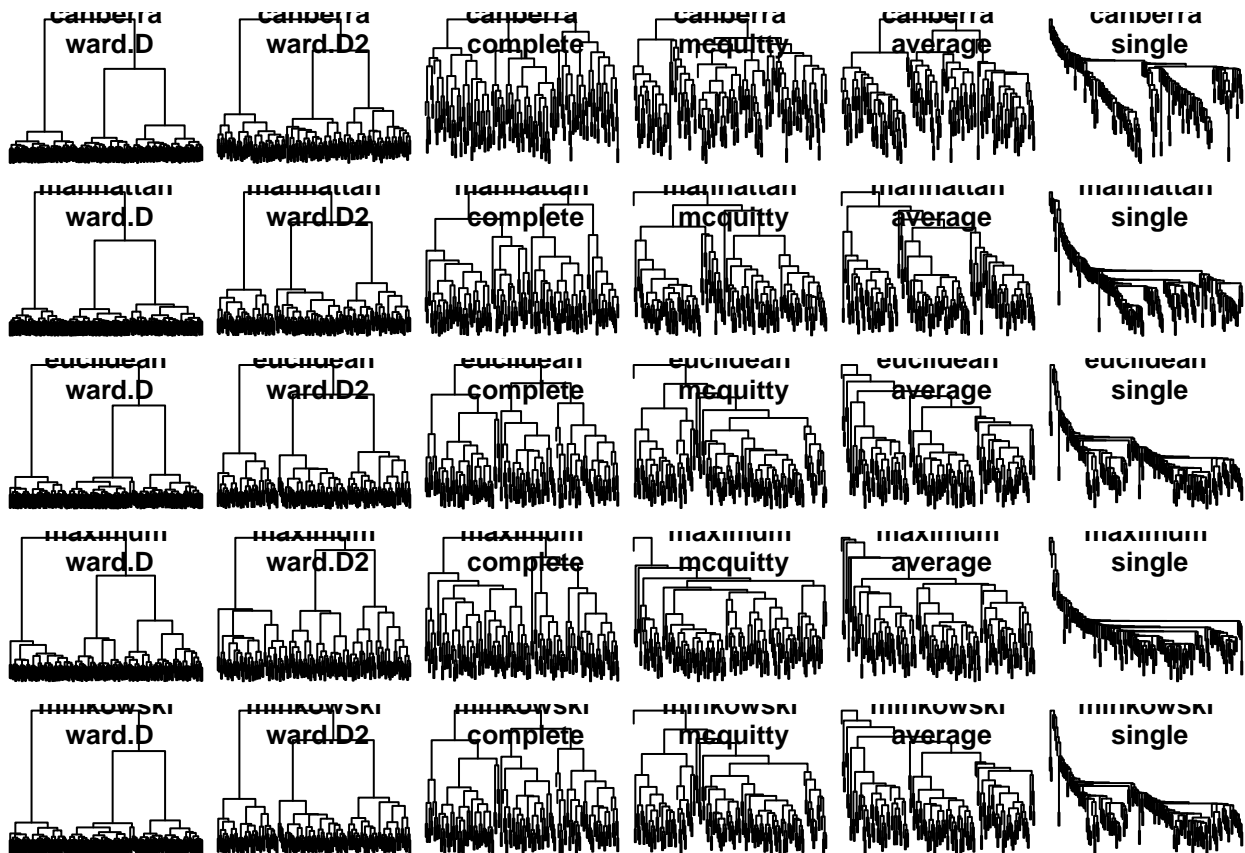
```

Below are different dendrograms for hierarchical clustering models having different clustering methods and distance metrics:

```


list_hclust <- list()
for(j in v_dist) for(k in v_hclust) list_hclust[[j]][[k]] <- hclust(
  d = list_dist[[j]],
  method = k
)
par(
  mfrow = c(length(v_dist), length(v_hclust)),
  mar = c(0,0,0,0),
  mai = c(0,0,0,0),
  oma = c(0,0,0,0)
)
for(j in v_dist) for(k in v_hclust) plot(
  x = list_hclust[[j]][[k]],
  labels = FALSE,
  axes = FALSE,
  main = paste("\n", j, "\n", k)
)

```

For comparison, the heights of these dendrograms should be on same scale. Thus the heights of the dendrograms are adjusted so that best two models can be chosen.


```
for(j in v_dist) for(k in v_hclust) list_hclust[[j]][[k]]$height <- rank(list_hclust[[j]][[k]]$height)
par(
  mfrow = c(length(v_dist),length(v_hclust)),
  mar = c(0,0,0,0),
  mai = c(0,0,0,0),
  oma = c(0,0,0,0)
)
for(j in v_dist) for(k in v_hclust) plot(
  x = list_hclust[[j]][[k]],
  labels = FALSE,
  axes = FALSE,
  main = paste("\n",j,"\n",k)
)
```



Based on the above matrix of dendrogram plots building a model that would be good for outlier detection:

```
plot(
  x = list_hclust[["manhattan"]][["single"]],
  main = "Manhattan Single Linkage",
  sub = ""
)
```

Manhattan Single Linkage



From the dendrogram created using manhattan single linkage hierarchical model, we can say that rows 159 & 160 are appearing distinct from the rest of the tree branches but are highly associated to each other. these can be potential outliers that need to be further investigated. Same is the case with observations in row 64 & 99.

Finding the clustering coefficient of this model:

```
cluster::coef.hclust(list_hclust[["manhattan"]][["single"]])
```


```
## [1] 0.5487526
```

The clustering coefficient for the hierarchical clustering model using “single” linkage method and distance metric created by “manhattan” method is low (around 0.548) which suggest this model to be good for outlier detection.

Based on the above matrix of dendrogram plots building a model that would be good for partitioning data into equal sized groups:


```
plot(
  x = list_hclust[["canberra"]][["ward.D"]],
  main = "Canberra Ward'D",
  sub = ""
)
```


Banner of `cluster::agnes(x = scale_M)`



Agglomerative Coefficient = 0.7

Dendrogram of cluster::agnes(x = scale_M)




From the dendrogram created by agnes() function we see that few of the records appear differently in the tree. the record of row 60 looks like an outlier in this dataset.

Thus we can say that agnes() function which uses agglomerative nesting algorithm for clustering is enabling us to detect outliers in the dataset.

Hierarchical clustering using diana() method


```
diana_M <- cluster::diana(scale_M)
plot(diana_M)
```

Banner of `cluster::diana(x = scale_M)`



Divisive Coefficient = 0.8

Dendrogram of cluster::diana(x = scale_M)



From the dendrogram created using Diana() function, we can say that the tree appears to be well partitioned. If we would cut the tree at a height between 8.7 to 8.9 (approx) we can get almost evenly partitioned clusters. Even the coefficient of clustering for the hierarchical model created using diana() is around 0.8 which suggests that well partitioned clusters can be formed of the dataset.

Hierarchical clustering using mona() method

To use mona() function there is need to convert data into binary format.

```
binary_M <- scale_M
for(j in 1:ncol(binary_M)) binary_M[,j] <- as.numeric(
  binary_M[,j] > median(binary_M[,j])
)
```

Now, as values of all variables are converted into 0's and 1's, monothetic analysis clustering of binary variables can be done (mona).

```
mona_M <- cluster::mona(binary_M)
print(mona_M)
```

```
## mona(x, ...) fit; x of dimension 178x13
## Order of objects:
## [1] 1 6 19 53 59 13 52 16 27 14 34 35 37 3 15 58 49 8
## [19] 7 4 20 40 43 46 29 54 5 17 18 31 36 56 25 72 2 28
## [37] 39 42 9 10 30 67 12 11 47 44 21 32 50 41 48 55 57 23
## [55] 33 22 24 38 66 45 51 96 82 101 64 81 94 95 100 98 99 85
```



```

## [73] 111 26 74 75 110 80 122 112 121 125 124 115 126 116 127 128 129 60
## [91] 77 102 131 63 73 68 105 117 70 86 118 61 79 71 69 62 155 65
## [109] 83 88 92 114 87 91 107 90 93 108 104 109 135 159 160 76 78 119
## [127] 120 136 168 171 172 145 164 84 133 138 148 161 158 132 152 140 163 134
## [145] 89 130 141 137 113 103 106 97 123 139 154 157 142 156 174 143 149 173
## [163] 178 144 147 165 166 146 176 162 167 169 175 177 150 151 170 153
## Variable used:
## [1] NULL NULL NULL
## [4] Hue Magnesium NULL
## [7] Proanthocyanins Magnesium Nonflavanoid_Phenols
## [10] Magnesium Color_Intensity Total_Phenols
## [13] Malic_Acid NULL NULL
## [16] OD280 Proanthocyanins Magnesium
## [19] Hue NULL NULL
## [22] NULL Proanthocyanins Nonflavanoid_Phenols
## [25] Color_Intensity Ash_Alcanity Color_Intensity
## [28] Malic_Acid Nonflavanoid_Phenols OD280
## [31] NULL Magnesium Proline
## [34] Ash Magnesium NULL
## [37] Malic_Acid Proanthocyanins NULL
## [40] NULL Proline Total_Phenols
## [43] Magnesium NULL Color_Intensity
## [46] Malic_Acid NULL NULL
## [49] Hue NULL NULL
## [52] NULL Color_Intensity Nonflavanoid_Phenols
## [55] Alcohol Malic_Acid OD280
## [58] Nonflavanoid_Phenols Ash Magnesium
## [61] Color_Intensity Total_Phenols Malic_Acid
## [64] Proline Proanthocyanins Malic_Acid
## [67] Proanthocyanins Nonflavanoid_Phenols OD280
## [70] NULL Hue Malic_Acid
## [73] Ash_Alcanity Malic_Acid Ash
## [76] Magnesium Proline Color_Intensity
## [79] Magnesium Proanthocyanins NULL
## [82] Ash Nonflavanoid_Phenols Malic_Acid
## [85] Proanthocyanins Hue Total_Phenols
## [88] Malic_Acid Flavanoids NULL
## [91] NULL Magnesium Alcohol
## [94] Ash_Alcanity OD280 NULL
## [97] Hue Magnesium Proanthocyanins
## [100] Ash_Alcanity Nonflavanoid_Phenols Proanthocyanins
## [103] Ash_Alcanity Alcohol Proline
## [106] Ash_Alcanity Ash Magnesium
## [109] OD280 Proanthocyanins Total_Phenols
## [112] Ash Ash_Alcanity OD280
## [115] Proanthocyanins Hue NULL
## [118] Ash_Alcanity Proanthocyanins Total_Phenols
## [121] Alcohol NULL Malic_Acid
## [124] Nonflavanoid_Phenols Magnesium Proanthocyanins
## [127] Color_Intensity NULL Proline
## [130] NULL Magnesium Nonflavanoid_Phenols
## [133] Ash_Alcanity Ash Nonflavanoid_Phenols
## [136] NULL NULL Proline
## [139] Magnesium NULL Nonflavanoid_Phenols

```

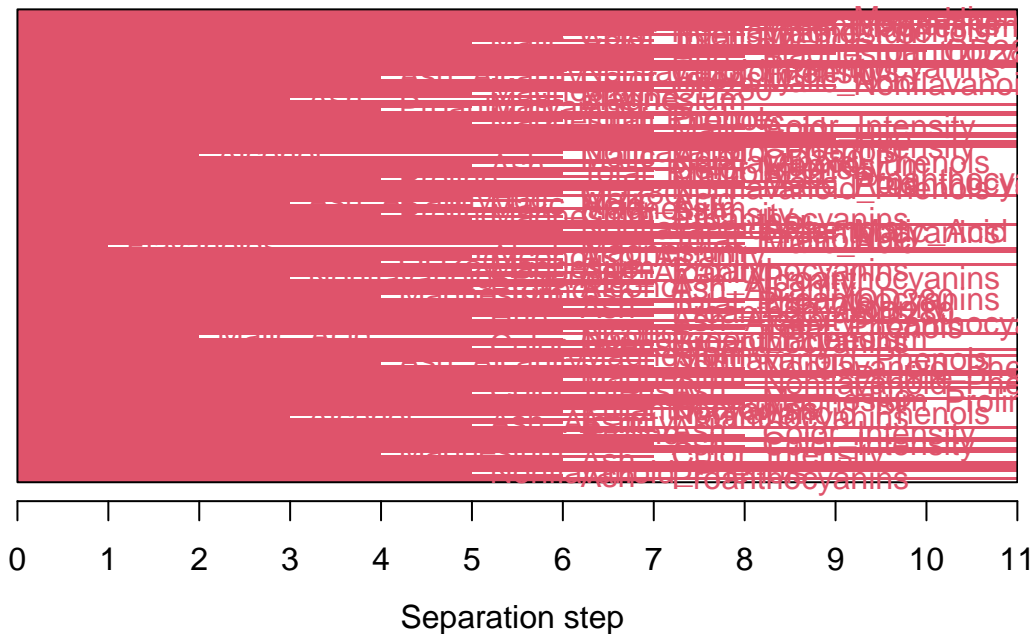
```

## [142] NULL          Ash          Color_Intensity
## [145] Hue             NULL          Proline
## [148] Magnesium       OD280         Ash
## [151] Proanthocyanins Nonflavanoid_Phenols Alcohol
## [154] Proanthocyanins Ash          Ash_Alcanity
## [157] NULL           Ash          Proline
## [160] Color_Intensity NULL          NULL
## [163] Ash            Color_Intensity NULL
## [166] NULL          Magnesium     Color_Intensity
## [169] Ash           NULL          NULL
## [172] NULL          NULL          Nonflavanoid_Phenols
## [175] Ash           NULL          Proanthocyanins
## Separation step:
## [1] 0 0 0 0 10 9 0 8 9 7 8 6 7 5 0 0 10 9 8 7 0 0 0 8 6 7
## [26] 4 6 8 9 7 0 5 6 3 6 0 5 4 0 0 7 6 5 0 8 7 0 0 9 0
## [51] 0 0 8 6 2 6 8 7 5 8 7 6 7 4 9 8 9 7 6 0 5 6 3 5 7
## [76] 6 4 6 5 7 0 8 6 9 8 9 7 8 1 0 0 6 5 6 4 0 6 5 7 6
## [101] 3 8 7 6 5 7 6 4 9 8 7 6 8 9 7 5 0 7 9 8 6 0 2 6 8
## [126] 7 5 0 7 0 6 7 4 7 8 0 0 9 6 0 8 0 7 5 9 0 10 8 7 8
## [151] 6 7 3 7 6 5 0 7 6 8 0 0 7 8 0 0 4 7 6 0 0 0 0 5 6
## [176] 0 7
##
## Available components:
## [1] "data"          "hasNA"         "order"         "variable"      "step"          "order.lab"
## [7] "call"

```

```
plot(mona_M)
```

Banner of cluster::mona(x = binary_M)



In mona(), each division is based on a single (well-chosen) variable unlike other hierarchical methods like agnes & hclust. But the original dataset used is not in binary format, thus results of this model would be biased.

Final Thoughts on Hierarchical Clustering Algorithm For Hierarchical clustering, below are the two best models:

- The hierarchical model build using “single” linkage method and distance metric created by “manhattan” method is having low (around 0.548) clustering coefficient which suggest this model to be good for outlier detection. Thus I would suggests that this as a best model for outlier detection.(As this is based on a comparative scale)
- The hierarchical model build using “ward.D” clustering method and distance metric created by “canberra” method is having high (around 0.701) clustering coefficient which suggests that this model to be good for partitioning the data into equal sized groups.(As this is based on a comparative scale). Thus I consider this to be my best model for partitioning the data into equal sized groups.

Some Interesting observations

- Models built by `agnes()` and `hclust()` with average linkage show that row number 60 is distinct from all other observations suggesting it to be an outlier in the data. However since the dendrograms of these models are not in the same height scale these are not chosen as final models.
- Model built using `diana()` has a clustering coefficient of 0.8 suggesting a good partition of data. By observing the dendrogram created by `diana()` we can say that the data can be clustered into a equal sized group of 3-4.